

## ▼ Exercise 1

### ▼ Exercise 1.1

```
for i in range(1,6):
    print("x"*i)
for y in range(1,6):
    print("x"*(5-y))
```

```
x
xx
xxx
xxxx
xxxxx
xxxx
xxx
xx
x
```

write a code that prints the following pattern. Try to use as few loops as possible.

```
X
X X
X X X
X X X X
X X X X X
X X X X
X X X
X X
X
```

Avoid using trivial solutions like:

```
print('X')
print('X X')
print('X X X')
print('X X X X')
print('X X X X X')
print('X X X X')
print('X X X')
print('X X')
print('X')
```

## ▼ Exercise 1.2

Write a script that will sum all numbers in the following string. You can split each number into single digits, so for example, you can consider 45 to be 4 and 5

```
sum = 0
input_str = "n45as29@#8ss6"
for i in input_str:
    if not i.isdigit():
        continue
    sum += int(i)
print( sum )
```

34

## ▼ Exercise 1.3

Write a script that will convert an arbitrary integer to a binary number (the number will be represented as a string with only `0` and `1`). Avoid using the `bin()` function or any other pythons default functions

```
def convert_to_binary(number: int) -> str:
    if number == 0:
        return "0"
    quotient = number // 2
    resto = number % 2
    binary = str(resto)
    while quotient != 0:
        resto = quotient % 2
        binary += str(resto)
        quotient = quotient // 2

    print(binary)
    return binary[::-1]
convert_to_binary(255)
```

```
11111111
'11111111'
```

## ▼ Exercise 1.4 - The Fibonacci Sequence:

The Fibonacci Sequence is a series of numbers. The following number is found by adding up the two numbers before it. The first two numbers are 0 and 1. For example, 0, 1, 1, 2, 3, 5, 8, 13. The following number in this series above is  $8 + 13 = 21$

Your task is to implement a function *fibonacci* that takes an integer as an input and returns a list that contains all Fibonacci numbers with values lower than the input integer

Example:

```
print(fibonacci(10))
```

```
[0, 1, 1, 2, 3, 5, 8]
```

```
def fibonacci(upper_threshold: int) -> list:
    result = []
    a, b = 0, 1

    while a < upper_threshold:
        result.append(a)
        a, b = b, a + b

    return result
```

```
fibonacci(22)
```

```
[0, 1, 1, 2, 3, 5, 8, 13, 21]
```

## ▼ Advanced

Try to implement the function using multiple approaches: *Iterative approach, Recursion, Memoization*

## ▼ Exercise 1.5. - Rock, Paper, Scissors game:

### ▼ basic

*Rock, Paper, Scissors* is a well-known and straightforward game. If you do not know the rules, google them. We will write a code for the Rock, Paper, Scissors game where the user plays against a random computer. The code can be written into one function *rock\_paper\_scissors*.

Notes:

- This implementation uses the **random** library to enable the computer to make a random choice.
- Keyword **Input** may be helpful. Check it
- After the function is run, the program will ask you about your movement, which you type in
- The game is case-insensitive for user input (e.g., "rock", "Rock", and "ROCK" are all valid).
- This script plays one round of the game.
- The script prints the result ('You lose', 'You win', 'It is a tie')

```
import random
def rock_paper_scissors() -> None:
    user_input = input("Enter your choice (rock, paper, scissors): ").lower()
    machine_choice = random.choice(["rock", "paper", "scissors"])
    print(f"Computer chose: {machine_choice}")
    if user_input not in ["rock", "paper", "scissors"]:
        print("Invalid input. Please enter rock, paper, or scissors.")
        return
    if user_input == "rock":
        if machine_choice == "scissors":
            print("You win!")
        elif machine_choice == "paper":
            print("You lose!")
        elif machine_choice == "rock":
            print("It's a tie!")
    if user_input == "scissors":
        if machine_choice == "paper":
            print("You win!")
        elif machine_choice == "rock":
            print("You lose!")
        elif machine_choice == "scissors":
            print("It's a tie!")
    if user_input == "paper":
        if machine_choice == "rock":
            print("You win!")
        elif machine_choice == "scissors":
            print("You lose!")
        elif machine_choice == "paper":
            print("It's a tie!")

rock_paper_scissors()
```

```
Enter your choice (rock, paper, scissors): paper
Computer chose: scissors
You lose!
```

## ▼ Advanced

Extend the Rock, Paper, Scissors game to be able to play  $n$  rounds.

## Notes:

- The scripts will contain two counters - user score and computer score
- After each run, the script prints the current scores and rounds.
- When the number of rounds is reached, the script prints the result ('You lose', 'You win', 'It is a tie')

```
import random
rounds=0
user_score=0
computer_score=0
victory_status= ""
i = 0
def rock_paper_scissors() -> None:
    global user_score, computer_score, i, victory_status
    rounds_to_play = int(input("Enter the number of rounds: "))
    while i < rounds_to_play:
        user_input = input("Enter your choice (rock, paper, scissors): ").lower()
        machine_choice = random.choice(["rock", "paper", "scissors"])
        print(f"Computer chose: {machine_choice}")
        if user_input not in ["rock", "paper", "scissors"]:
            print("Invalid input. Please enter rock, paper, or scissors.")
            continue # Continue to the next iteration instead of returning

        if user_input == "rock":
            if machine_choice == "scissors":
                victory_status="You win!"
            elif machine_choice == "paper":
                victory_status="You lose!"
            elif machine_choice == "rock":
                victory_status="It's a tie!"
        elif user_input == "scissors":
            if machine_choice == "paper":
                victory_status="You win!"
            elif machine_choice == "rock":
                victory_status="You lose!"
            elif machine_choice == "scissors":
                victory_status="It's a tie!"
        elif user_input == "paper": # Use elif
            if machine_choice == "rock":
                victory_status="You win!"
            elif machine_choice == "scissors":
                victory_status="You lose!"
            elif machine_choice == "paper":
                victory_status="It's a tie!"

    print(victory_status)

    if victory_status=="You win!":
        user_score+=1
    elif victory_status=="You lose!":
        computer_score+=1
```

```
i += 1 # Increment the round counter
print(f"Current Scores: Player {user_score} - Computer {computer_score}")

print("\n--- Final Results ---")
if user_score > computer_score:
    print("You win the game!")
elif computer_score > user_score:
    print("The computer wins the game!")
else:
    print("The game is a tie!")

rock_paper_scissors()

Enter the number of rounds: 3
Enter your choice (rock, paper, scissors): paper
Computer chose: paper
It's a tie!
Current Scores: Player 0 - Computer 0 (Round 1/3)
Enter your choice (rock, paper, scissors): paper
Computer chose: rock
You win!
Current Scores: Player 1 - Computer 0 (Round 2/3)
Enter your choice (rock, paper, scissors): paper
Computer chose: paper
It's a tie!
Current Scores: Player 1 - Computer 0 (Round 3/3)

--- Final Results ---
You win the game!
```

## ▼ Exercise 2

The purpose of this exercise is to practise working with NumPy library

### ▼ Exercise 2.1

In this exercise, you will work with NumPy arrays and learn how to efficiently modify and process them using both loops and vectorized operations. Your task is to:

1. Implement the function *create\_array\_nxn* that generates an  $n \times n$  NumPy array filled with numbers from  $n^2 - 1$  down to 0
2. Using a loop-based approach (for-loop, while-loop), implement the function *apply\_threshold\_loop* that takes any NumPy array and replaces all numbers below a user-defined threshold with 0.
3. Implement function *apply\_threshold\_vectorized* that performs the same operation as *apply\_threshold\_loop*, but without loops, using NumPy vectorized

operations

4. Implement function *compare\_performance* that compares the execution time of functions *apply\_threshold\_vectorized* and *apply\_threshold\_loop* with **time** library and print the results

### Additional Notes:

- Ensure n is a positive integer ( $n > 0$ ).
- Use the time library to measure execution times for performance comparison.
- The function *compare\_performance(n, threshold)* should:
  - Generate an  $n \times n$  array.
  - Apply both loop-based and vectorized thresholding.
  - Print execution times.

```
import numpy as np
import time
```

```
def create_array_nxn(n: int) -> np.ndarray:
    if n <= 0:
        raise ValueError("n must be a positive integer")

    # Create a 1D array with numbers from n*n - 1 down to 0
    flat_array = np.arange(n**2 - 1, -1, -1)

    # Reshape the 1D array into an n x n 2D array
    my_array = flat_array.reshape((n, n))

    return my_array
```

```
create_array_nxn(5)
```

```
array([[24, 23, 22, 21, 20],
       [19, 18, 17, 16, 15],
       [14, 13, 12, 11, 10],
       [ 9,  8,  7,  6,  5],
       [ 4,  3,  2,  1,  0]])
```

```
def apply_threshold_loop(arr: np.ndarray, threshold: int) -> np.ndarray:
    # Create a copy to avoid modifying the original array
    modified_arr = arr.copy()

    # Iterate through the array using nested loops
    for i in range(modified_arr.shape[0]):
        for j in range(modified_arr.shape[1]):
            if modified_arr[i, j] < threshold:
                modified_arr[i, j] = 0

    return modified_arr
```

```

def apply_threshold_vectorized(arr: np.ndarray, threshold: int) -> np.ndarray
    modified_arr = arr.copy()
    modified_arr[modified_arr < threshold] = 0
    return modified_arr

```

```

def compare_performance(n: int, threshold: int) -> None:
    # Generate the n x n array
    arr = create_array_nxn(n)

    # Measure performance of loop-based approach
    start_time_loop = time.time()
    apply_threshold_loop(arr, threshold)
    end_time_loop = time.time()
    loop_time = end_time_loop - start_time_loop

    # Measure performance of vectorized approach
    start_time_vectorized = time.time()
    apply_threshold_vectorized(arr, threshold)
    end_time_vectorized = time.time()
    vectorized_time = end_time_vectorized - start_time_vectorized

    # Print results
    print(f"Performance comparison for n={n}, threshold={threshold}:")
    print(f"  Loop-based approach: {loop_time:.6f} seconds")
    print(f"  Vectorized approach: {vectorized_time:.6f} seconds")

    # Optionally, print the ratio for clearer comparison
    if vectorized_time > 0:
        print(f"  Vectorized is {loop_time / vectorized_time:.2f}x faster than loop-based")
    else:
        print("  Vectorized time was zero, cannot calculate speedup ratio.")

```

## Exercise 2.2 - Digi display

basic

The task will be to create a simulation of a Digi display that will be able to display an arbitrary integer

\*hint: to show the image, use the library `matplotlib.pyplot`. Also, you may find function `np.concatenate` helpful

```

nums = {
    "1": np.array([[0, 1, 1], [1, 0, 1], [0, 0, 1], [0, 0, 1], [0, 0, 1]])
    "2": np.array([[1, 1, 1], [0, 0, 1], [1, 1, 1], [1, 0, 0], [1, 1, 1]])
    "3": np.array([[1, 1, 1], [0, 0, 1], [1, 1, 1], [0, 0, 1], [1, 1, 1]])
    "4": np.array([[1, 0, 1], [1, 0, 1], [1, 1, 1], [0, 0, 1], [0, 0, 1]])
    "5": np.array([[1, 1, 1], [1, 0, 0], [1, 1, 1], [0, 0, 1], [1, 1, 1]])
    "6": np.array([[1, 1, 1], [1, 0, 0], [1, 1, 1], [1, 0, 1], [1, 1, 1]])
    "7": np.array([[1, 1, 1], [0, 0, 1], [0, 0, 1], [0, 0, 1], [0, 0, 1]])
    "8": np.array([[1, 1, 1], [1, 0, 1], [1, 1, 1], [1, 0, 1], [1, 1, 1]])
}

```

```
"9": np.array([[1, 1, 1], [1, 0, 1], [1, 1, 1], [0, 0, 1], [1, 1, 1]])
"0": np.array([[1, 1, 1], [1, 0, 1], [1, 0, 1], [1, 0, 1], [1, 1, 1]])
}
```

```
import matplotlib.pyplot as plt
def show_in_digi(input_integer: int) -> None:
    s = str(input_integer)
    digit_images = []
    for digit_char in s:
        if digit_char in numbs:
            digit_images.append(numbs[digit_char])
        else:
            # Handle cases where character is not a digit (e.g., negative si
            # For now, let's assume valid digits. Could add a blank or error
            pass # Or raise an error, or use a placeholder

    if not digit_images:
        print("No valid digits to display.")
        return

    # Concatenate digit images horizontally
    # Add a small separator (e.g., a column of zeros) between digits for bet
    combined_image = digit_images[0]
    for i in range(1, len(digit_images)):
        # Add a blank column as separator
        separator = np.zeros((digit_images[0].shape[0], 1))
        combined_image = np.concatenate((combined_image, separator, digit_im

    plt.imshow(combined_image, cmap='binary') # Use 'binary' colormap for bl
    plt.axis('off') # Hide axes
    plt.show()

show_in_digi(5289)
```



```
show_in_digi(5289)
```

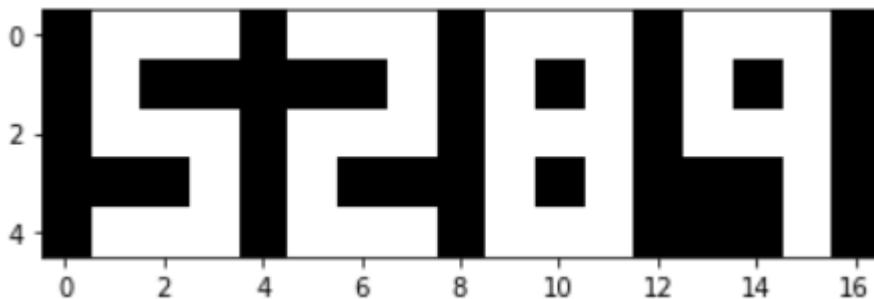
```
NameError Traceback (most recent call last)
/tmp/ipython-input-245880815.py in <cell line: 0>()
----> 1 show_in_digi(5289)

/tmp/ipython-input-2794560182.py in show_in_digi(input_integer)
    4     digit_images = []
    5     for digit_char in s:
----> 6         if digit_char in numbs:
    7             digit_images.append(numbs[digit_char])
    8         else:

NameError: name 'numbs' is not defined
```

Étapes suivantes : [Expliquer l'erreur](#)

`show_in_digi(5289)` will show:



*hint2: you may find the following dict usefull*

```
numbs = {
    "1": np.array([[0, 1, 1], [1, 0, 1], [0, 0, 1], [0, 0, 1], [0, 0, 1]]),
    "2": np.array([[1, 1, 1], [0, 0, 1], [1, 1, 1], [1, 0, 0], [1, 1, 1]]),
    "3": np.array([[1, 1, 1], [0, 0, 1], [1, 1, 1], [0, 0, 1], [1, 1, 1]]),
    "4": np.array([[1, 0, 1], [1, 0, 1], [1, 1, 1], [0, 0, 1], [0, 0, 1]]),
    "5": np.array([[1, 1, 1], [1, 0, 0], [1, 1, 1], [0, 0, 1], [1, 1, 1]]),
    "6": np.array([[1, 1, 1], [1, 0, 0], [1, 1, 1], [1, 0, 1], [1, 1, 1]]),
    "7": np.array([[1, 1, 1], [0, 0, 1], [0, 0, 1], [0, 0, 1], [0, 0, 1]]),
    "8": np.array([[1, 1, 1], [1, 0, 1], [1, 1, 1], [1, 0, 1], [1, 1, 1]]),
    "9": np.array([[1, 1, 1], [1, 0, 1], [1, 1, 1], [0, 0, 1], [1, 1, 1]]),
    "0": np.array([[1, 1, 1], [1, 0, 1], [1, 0, 1], [1, 0, 1], [1, 1, 1]])
}
```

## ▼ advanced

Extend your `show_in_digi` function to be able to display an arbitrary float number and negative numbers

```
import numpy as np
import matplotlib.pyplot as plt
from typing import Union

# Extended nums dictionary to include negative sign and decimal point
nums = {
    "1": np.array([[0, 1, 1], [1, 0, 1], [0, 0, 1], [0, 0, 1], [0, 0, 1]]),
    "2": np.array([[1, 1, 1], [0, 0, 1], [1, 1, 1], [1, 0, 0], [1, 1, 1]]),
    "3": np.array([[1, 1, 1], [0, 0, 1], [1, 1, 1], [0, 0, 1], [1, 1, 1]]),
    "4": np.array([[1, 0, 1], [1, 0, 1], [1, 1, 1], [0, 0, 1], [0, 0, 1]]),
    "5": np.array([[1, 1, 1], [1, 0, 0], [1, 1, 1], [0, 0, 1], [1, 1, 1]]),
    "6": np.array([[1, 1, 1], [1, 0, 0], [1, 1, 1], [1, 0, 1], [1, 1, 1]]),
    "7": np.array([[1, 1, 1], [0, 0, 1], [0, 0, 1], [0, 0, 1], [0, 0, 1]]),
    "8": np.array([[1, 1, 1], [1, 0, 1], [1, 1, 1], [1, 0, 1], [1, 1, 1]]),
    "9": np.array([[1, 1, 1], [1, 0, 1], [1, 1, 1], [0, 0, 1], [1, 1, 1]]),
    "0": np.array([[1, 1, 1], [1, 0, 1], [1, 0, 1], [1, 0, 1], [1, 1, 1]]),
    "-": np.array([[0,0,0], [0,0,0], [1,1,1], [0,0,0], [0,0,0]]), # Negative
    ".": np.array([[0,0,0], [0,0,0], [0,0,0], [0,0,0], [0,1,0]]), # Decimal
}

def show_in_digi(input_number: Union[int, float]) -> None:
    s = str(input_number)
    digit_images = []
    for digit_char in s:
        if digit_char in nums:
            digit_images.append(nums[digit_char])
        else:
            print(f"Warning: Character '{digit_char}' not supported and will continue")

    if not digit_images:
        print("No valid digits or symbols to display.")
        return

    # Concatenate digit images horizontally
    # Add a small separator (e.g., a column of zeros) between digits for better readability
    # The separator width can be adjusted. Here, 1 column of zeros.
    separator_width = 1
    combined_image = digit_images[0]
    for i in range(1, len(digit_images)):
        separator = np.zeros((digit_images[0].shape[0], separator_width))
        combined_image = np.concatenate((combined_image, separator, digit_im

    plt.imshow(combined_image, cmap='binary') # Use 'binary' colormap for better contrast
    plt.axis('off') # Hide axes
    plt.show()

# Test with float and negative numbers
print("Displaying 12.34:")
show_in_digi(12.34)

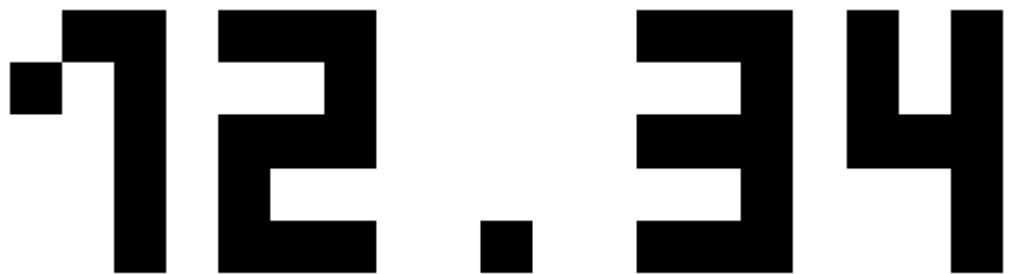
print("\nDisplaying -567:")
show_in_digi(-567)

print("\nDisplaying -0.89:")
show_in_digi(-0.89)
```

```
show_in_digi(-0.89)

print("\nDisplaying 0:")
show_in_digi(0)
```

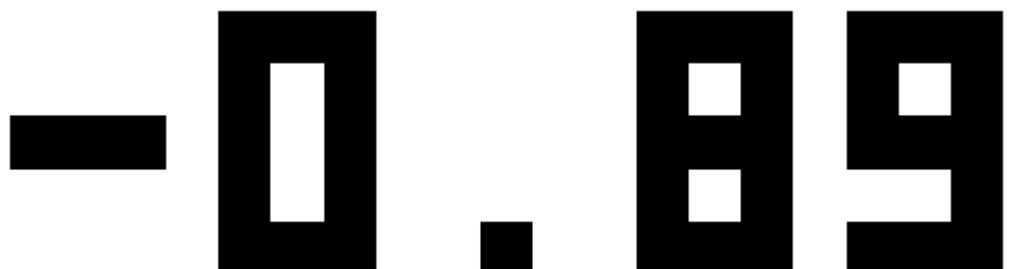
Displaying 12.34:



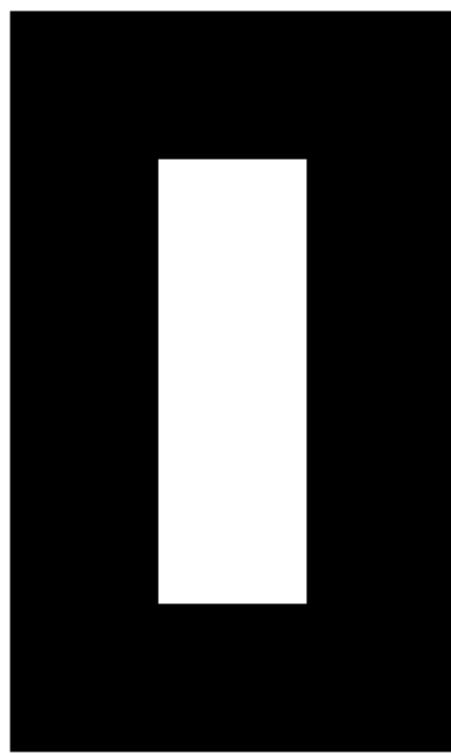
Displaying -567:



Displaying -0.89:



Displaying 0:



## Exercise 3 - Playing with California Housing Dataset

The purpose of this exercise is to learn basic operations of **pandas** library. Use *california\_housing\_dataset* you can find in google colab to do following operations:

- check what *dataset.describe()* does
- display all rows where *total\_bedrooms* column is bigger than 310
- drop also first and last row
- save a mean of values in *households* columns to variable, using *matplotlib* display number of households in the graph (x-axis would be ID, y-axis - numbers of households) as dots with the same color, display also mean as a single line in the graph with different color as dots is
- check if any of the columns contain NaN values, if yes, replace the NaN values with the arithmetic mean of the whole dataset
- display *lat* and *long* into plot, *lat* would be on x axis and *long* on y axis
- choose two arbitrary columns and normalize all values inside using *min-max normalization*:
- Create a Correlation Matrix from all columns

Commencez à coder ou à générer avec l'IA.

## Task

Perform the following operations on the `california_housing_train.csv` dataset: load it into a pandas DataFrame, describe the dataset to get statistical summaries, filter the DataFrame to include only rows where 'total\_bedrooms' is greater than 310, and then drop the first and last rows from this filtered DataFrame. Next, calculate the mean of the 'households' column and create a scatter plot showing 'households' against the DataFrame index, adding a horizontal line for the calculated mean. Check for any NaN values in the DataFrame; if found, replace them with the mean of all numeric values in the DataFrame. Then, create a scatter plot of 'latitude' (x-axis)

against 'longitude' (y-axis). Choose two arbitrary numerical columns (e.g., 'median\_income' and 'housing\_median\_age') and apply min-max normalization to them. Finally, generate and display a correlation matrix for all numerical columns in the DataFrame and summarize the basic pandas operations and key observations from the analysis.

## ▼ Load Dataset

### Subtask:

Load the 'california\_housing\_train.csv' dataset into a pandas DataFrame.

**Reasoning:** First, I'll import the pandas library, which is essential for data manipulation in Python, especially for handling DataFrames. Then, I'll use `pd.read_csv` to load the specified CSV file into a DataFrame named `df` and display its head to confirm the loading process.

```
import pandas as pd

df = pd.read_csv('/content/sample_data/california_housing_train.csv')
print("Dataset loaded successfully. Displaying the first 5 rows:")
print(df.head())
```

```
Dataset loaded successfully. Displaying the first 5 rows:
   longitude  latitude  housing_median_age  total_rooms  total_bedrooms \
0      -114.31      34.19             15.0       5612.0            1283.0
1      -114.47      34.40             19.0       7650.0            1901.0
2      -114.56      33.69             17.0       720.0             174.0
3      -114.57      33.64             14.0      1501.0            337.0
4      -114.57      33.57             20.0      1454.0            326.0

   population  households  median_income  median_house_value
0      1015.0        472.0        1.4936          66900.0
1      1129.0        463.0        1.8200          80100.0
2       333.0        117.0        1.6509          85700.0
3       515.0        226.0        3.1917          73400.0
4       624.0        262.0        1.9250          65500.0
```

## ▼ Describe Dataset

### Subtask:

Utilize the `.describe()` method on the loaded DataFrame to get a statistical summary of the numerical columns, including count, mean, standard deviation, min, max, and quartile values.

**Reasoning:** To get a statistical summary of the numerical columns in the DataFrame, I will use the `.describe()` method on the `df` DataFrame and print the result.

```
print("Descriptive statistics for the dataset:")
print(df.describe())
```

Descriptive statistics for the dataset:

	longitude	latitude	housing_median_age	total_rooms	\
count	17000.000000	17000.000000	17000.000000	17000.000000	
mean	-119.562108	35.625225	28.589353	2643.664412	
std	2.005166	2.137340	12.586937	2179.947071	
min	-124.350000	32.540000	1.000000	2.000000	
25%	-121.790000	33.930000	18.000000	1462.000000	
50%	-118.490000	34.250000	29.000000	2127.000000	
75%	-118.000000	37.720000	37.000000	3151.250000	
max	-114.310000	41.950000	52.000000	37937.000000	
	total_bedrooms	population	households	median_income	\
count	17000.000000	17000.000000	17000.000000	17000.000000	
mean	539.410824	1429.573941	501.221941	3.883578	
std	421.499452	1147.852959	384.520841	1.908157	
min	1.000000	3.000000	1.000000	0.499900	
25%	297.000000	790.000000	282.000000	2.566375	
50%	434.000000	1167.000000	409.000000	3.544600	
75%	648.250000	1721.000000	605.250000	4.767000	
max	6445.000000	35682.000000	6082.000000	15.000100	
	median_house_value				
count	17000.000000				
mean	207300.912353				
std	115983.764387				
min	14999.000000				
25%	119400.000000				
50%	180400.000000				
75%	265000.000000				
max	500001.000000				

**Reasoning:** The next step is to filter the DataFrame `df` to include only rows where the 'total\_bedrooms' column is greater than 310, as specified in the instructions. I will then display the head of the new filtered DataFrame to confirm the operation.

```
df_filtered = df[df['total_bedrooms'] > 310]
print("DataFrame filtered for 'total_bedrooms' > 310. Displaying the first 5")
print(df_filtered.head())
```

DataFrame filtered for 'total\_bedrooms' > 310. Displaying the first 5 rows:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-114.31	34.19	15.0	5612.0	1283.0	
1	-114.47	34.40	19.0	7650.0	1901.0	
3	-114.57	33.64	14.0	1501.0	337.0	
4	-114.57	33.57	20.0	1454.0	326.0	
6	-114.58	33.61	25.0	2907.0	680.0	

population households median\_income median\_house\_value

0	1015.0	472.0	1.4936	66900.0
1	1129.0	463.0	1.8200	80100.0
3	515.0	226.0	3.1917	73400.0
4	624.0	262.0	1.9250	65500.0
6	1841.0	633.0	2.6768	82400.0

**Reasoning:** To drop the first and last rows from the `df_filtered` DataFrame, I will use `iloc` with slicing `[1:-1]` to select all rows except the first and the last. Then, I will display the head and tail of the modified DataFrame to confirm the operation.

```
df_modified = df_filtered.iloc[1:-1]
print("DataFrame after dropping the first and last rows. Displaying the first 5 rows:")
print(df_modified.head())
print("\nDisplaying the last 5 rows:")
print(df_modified.tail())
```

DataFrame after dropping the first and last rows. Displaying the first 5 rows

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms
1	-114.47	34.40	19.0	7650.0	1901.0
3	-114.57	33.64	14.0	1501.0	337.0
4	-114.57	33.57	20.0	1454.0	326.0
6	-114.58	33.61	25.0	2907.0	680.0
8	-114.59	33.61	34.0	4789.0	1175.0

	population	households	median_income	median_house_value
1	1129.0	463.0	1.8200	80100.0
3	515.0	226.0	3.1917	73400.0
4	624.0	262.0	1.9250	65500.0
6	1841.0	633.0	2.6768	82400.0
8	3134.0	1056.0	2.1782	58400.0

Displaying the last 5 rows:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms
16993	-124.23	40.54	52.0	2694.0	453.0
16994	-124.25	40.28	32.0	1430.0	419.0
16995	-124.26	40.58	52.0	2217.0	394.0
16996	-124.27	40.69	36.0	2349.0	528.0
16997	-124.30	41.84	17.0	2677.0	531.0

	population	households	median_income	median_house_value
16993	1152.0	435.0	3.0806	106700.0
16994	434.0	187.0	1.9417	76100.0
16995	907.0	369.0	2.3571	111400.0
16996	1194.0	465.0	2.5179	79000.0
16997	1244.0	456.0	3.0313	103600.0

**Reasoning:** Now, I will calculate the mean of the 'households' column from the `df_modified` DataFrame. Then, I will create a scatter plot of 'households' against the DataFrame index, adding a horizontal line to represent the calculated mean, and ensure the plot has appropriate labels and a title for better visualization.

```

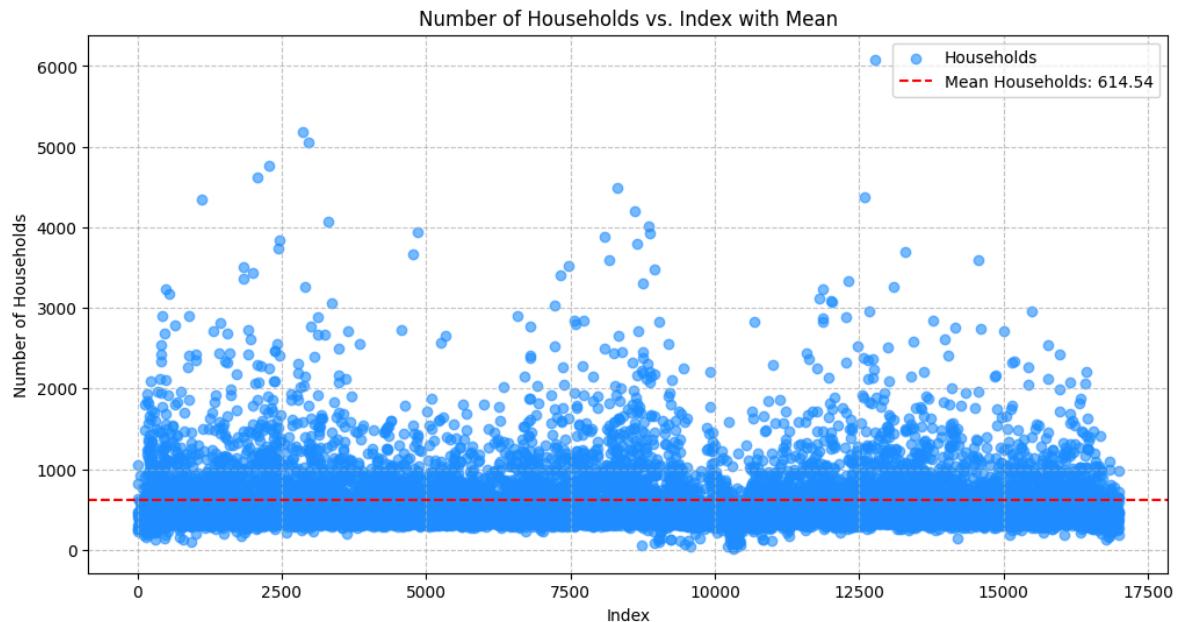
import matplotlib.pyplot as plt

mean_households = df_modified['households'].mean()
print(f"Mean of 'households' column: {mean_households:.2f}")

plt.figure(figsize=(12, 6))
plt.scatter(df_modified.index, df_modified['households'], alpha=0.6, color='blue')
plt.axhline(y=mean_households, color='red', linestyle='--', label=f'Mean Households: {mean_households:.2f}')
plt.xlabel('Index')
plt.ylabel('Number of Households')
plt.title('Number of Households vs. Index with Mean')
plt.legend()
plt.grid(True, linestyle='--', alpha=0.7)
plt.show()

```

Mean of 'households' column: 614.54



**Reasoning:** Now I will check for any NaN values in the `df_modified` DataFrame. If NaN values are found, I will replace them with the arithmetic mean of all numerical values in the DataFrame to handle missing data.

```

print("Checking for NaN values:")
print(df_modified.isnull().sum())

```

```
if df_modified.isnull().sum().any():
    # Calculate the mean of all numerical values in the DataFrame
    # First, select only numerical columns to calculate the mean
    numerical_cols = df_modified.select_dtypes(include=np.number).columns
    overall_mean = df_modified[numerical_cols].stack().mean()

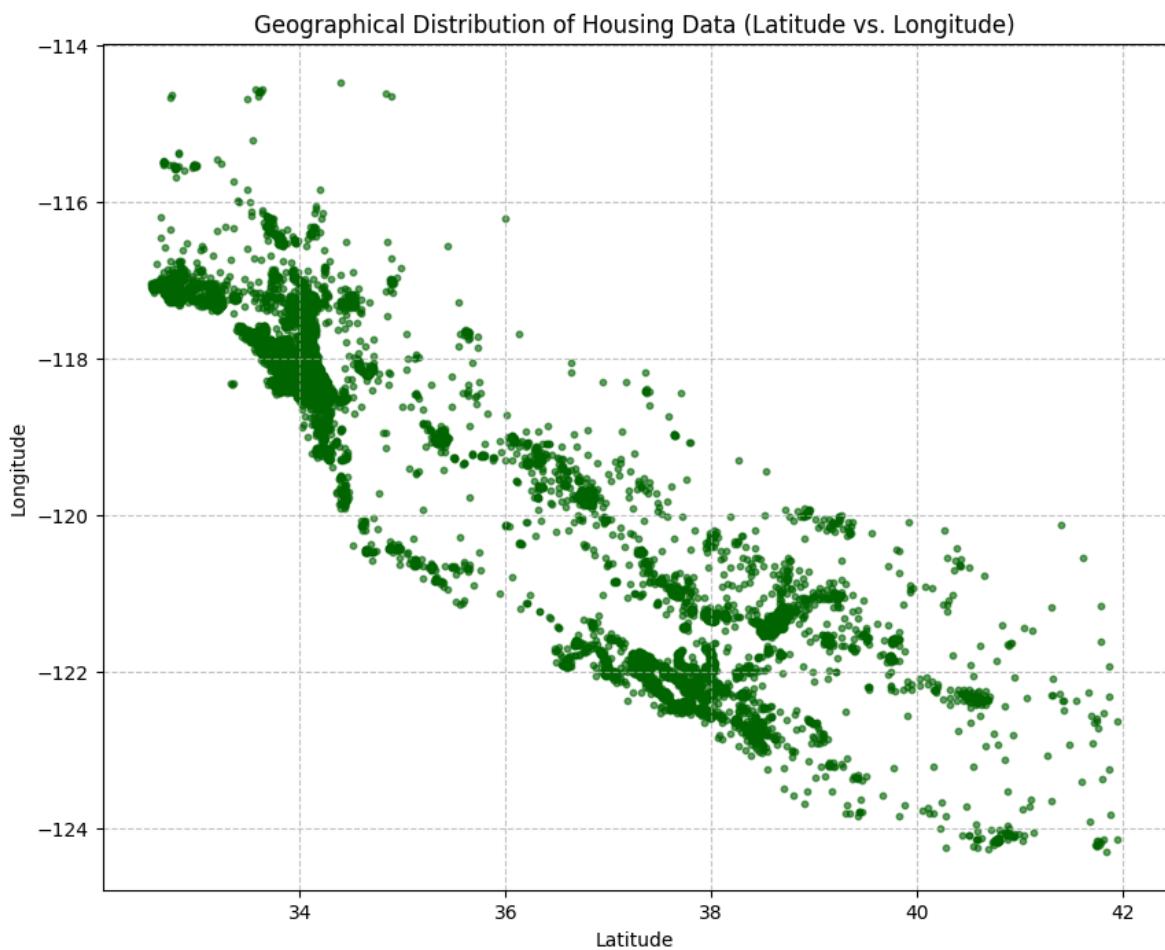
    print(f"\nNaN values found. Replacing with overall mean: {overall_mean:.2f}")
    df_modified.fillna(overall_mean, inplace=True)
    print("NaN values after replacement:")
    print(df_modified.isnull().sum())
else:
    print("\nNo NaN values found in the DataFrame.")

Checking for NaN values:
longitude          0
latitude           0
housing_median_age 0
total_rooms         0
total_bedrooms      0
population          0
households          0
median_income        0
median_house_value   0
dtype: int64

No NaN values found in the DataFrame.
```

**Reasoning:** Now, I will create a scatter plot to visualize the spatial distribution of the housing data, with 'latitude' on the x-axis and 'longitude' on the y-axis, as requested.

```
plt.figure(figsize=(10, 8))
plt.scatter(df_modified['longitude'], df_modified['latitude'], alpha=0.6, color='blue')
plt.xlabel('Latitude')
plt.ylabel('Longitude')
plt.title('Geographical Distribution of Housing Data (Latitude vs. Longitude)')
plt.grid(True, linestyle='--', alpha=0.7)
plt.show()
```



**Reasoning:** Now, I will choose two arbitrary numerical columns from the `df_modified` DataFrame, 'median\_income' and 'housing\_median\_age', and apply min-max normalization to them to scale their values between 0 and 1. I will then display the head of the DataFrame to show the normalized columns.

```
selected_cols = ['median_income', 'housing_median_age']

for col in selected_cols:
    min_val = df_modified[col].min()
    max_val = df_modified[col].max()
    df_modified[f'{col}_normalized'] = (df_modified[col] - min_val) / (max_v
```

```
print("DataFrame after Min-Max normalization for 'median_income' and 'housin
print(df_modified[['median_income', 'median_income_normalized', 'housing_med
```

```
DataFrame after Min-Max normalization for 'median_income' and 'housing_median
median_income  median_income_normalized  housing_median_age \
1           1.8200                 0.079769      19.0
3           3.1917                 0.175541      14.0
4           1.9250                 0.087100      20.0
6           2.6768                 0.139591      25.0
8           2.1782                 0.104778      34.0
```

```
housing_median_age_normalized
1                  0.352941
3                  0.254902
4                  0.372549
6                  0.470588
8                  0.647059
```

```
/tmp/ipython-input-2954841273.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/api/index.html#pandas.DataFrame.set\\_value](https://pandas.pydata.org/pandas-docs/stable/api/index.html#pandas.DataFrame.set_value)  
`df_modified[f'{col}_normalized'] = (df_modified[col] - min_val) / (max_val - min_val)`

```
/tmp/ipython-input-2954841273.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/api/index.html#pandas.DataFrame.set\\_value](https://pandas.pydata.org/pandas-docs/stable/api/index.html#pandas.DataFrame.set_value)  
`df_modified[f'{col}_normalized'] = (df_modified[col] - min_val) / (max_val - min_val)`

**Reasoning:** The previous code produced a `SettingWithCopyWarning`. This warning indicates that operations are being performed on a slice of a DataFrame, which might not modify the original DataFrame as intended. To resolve this, I will explicitly create a deep copy of the DataFrame `df_filtered` when creating `df_modified` to ensure independent modification. Then, I will perform min-max normalization on the selected columns and display the results.

```
df_modified = df_filtered.iloc[1:-1].copy() # Ensure df_modified is a deep copy

selected_cols = ['median_income', 'housing_median_age']

for col in selected_cols:
    min_val = df_modified[col].min()
    max_val = df_modified[col].max()
    df_modified[f'{col}_normalized'] = (df_modified[col] - min_val) / (max_val - min_val)
```

```

print("DataFrame after Min-Max normalization for 'median_income' and 'housing_median_age'")
print(df_modified[['median_income', 'median_income_normalized', 'housing_median_age_normalized']])

DataFrame after Min-Max normalization for 'median_income' and 'housing_median_age'
   median_income  median_income_normalized  housing_median_age_normalized \
1            1.8200                  0.079769                      19.0
3            3.1917                  0.175541                      14.0
4            1.9250                  0.087100                      20.0
6            2.6768                  0.139591                      25.0
8            2.1782                  0.104778                      34.0

   housing_median_age_normalized \
1                  0.352941
3                  0.254902
4                  0.372549
6                  0.470588
8                  0.647059

```

**Reasoning:** Now, I will generate a correlation matrix for all numerical columns in the `df_modified` DataFrame to understand the relationships between different features. I will then display this matrix.

```

print("Correlation Matrix for all numerical columns:")
correlation_matrix = df_modified.corr()
print(correlation_matrix)

```

```

Correlation Matrix for all numerical columns:
   longitude  latitude  housing_median_age \
longitude      1.000000 -0.926237          -0.121227
latitude       -0.926237  1.000000           0.024286
housing_median_age -0.121227  0.024286          1.000000
total_rooms     0.041182 -0.023179          -0.375131
total_bedrooms   0.070216 -0.060194          -0.324652
population      0.096686 -0.100238          -0.296564
households      0.052823 -0.063723          -0.304374
median_income    -0.017161 -0.075634          -0.140311
median_house_value -0.053489 -0.139220           0.132600
median_income_normalized -0.017161 -0.075634          -0.140311
housing_median_age_normalized -0.121227  0.024286          1.000000

   total_rooms  total_bedrooms  population \
longitude        0.041182      0.070216  0.096686
latitude        -0.023179     -0.060194 -0.100238
housing_median_age -0.375131     -0.324652 -0.296564
total_rooms       1.000000      0.911944  0.836176
total_bedrooms      0.911944      1.000000  0.858352
population        0.836176      0.858352  1.000000
households        0.899982      0.975801  0.893430
median_income      0.260194      0.006390  0.025688
median_house_value  0.132169      0.039259 -0.046994
median_income_normalized  0.260194      0.006390  0.025688
housing_median_age_normalized -0.375131     -0.324652 -0.296564

   households  median_income  median_house_value \
longitude        0.052823     -0.017161          -0.053489

```

latitude	-0.063723	-0.075634	-0.139226
housing_median_age	-0.304374	-0.140311	0.132606
total_rooms	0.899982	0.260194	0.132169
total_bedrooms	0.975801	0.006390	0.039256
population	0.893430	0.025688	-0.046994
households	1.000000	0.031419	0.057071
median_income	0.031419	1.000000	0.678163
median_house_value	0.057077	0.678163	1.000000
median_income_normalized	0.031419	1.000000	0.678163
housing_median_age_normalized	-0.304374	-0.140311	0.132606
 median_income_normalized \			
longitude	-0.017161		
latitude	-0.075634		
housing_median_age	-0.140311		
total_rooms	0.260194		
total_bedrooms	0.006390		
population	0.025688		
households	0.031419		
median_income	1.000000		
median_house_value	0.678163		
median_income_normalized	1.000000		
housing_median_age_normalized	-0.140311		
 housing_median_age_normalized			
longitude	-0.121227		
latitude	0.024286		
housing_median_age	1.000000		

## Summary of Operations and Observations

In this exercise, we performed several key data analysis operations using the pandas library on the `california_housing_train.csv` dataset. The steps included:

- 1. Loading the Dataset:** We successfully loaded the `california_housing_train.csv` into a pandas DataFrame named `df`.
- 2. Descriptive Statistics:** We utilized the `.describe()` method to obtain a statistical summary of the numerical columns. This provided insights into the count, mean, standard deviation, min, max, and quartile values for each feature, offering a quick overview of the data distribution.
- 3. Filtering Data:** We filtered the DataFrame to include only rows where the 'total\_bedrooms' column was greater than 310, creating `df_filtered`. This allowed us to focus on properties with a significant number of bedrooms.
- 4. Dropping Rows:** From the filtered DataFrame, we dropped the first and last rows, resulting in `df_modified`. This operation demonstrated how to selectively remove data entries from the beginning and end of a DataFrame.
- 5. Calculating Mean and Visualization:** We calculated the mean of the 'households' column and then created a scatter plot. The plot visualized the

distribution of households across the dataset, with a horizontal line indicating the mean, providing a clear visual reference.

6. **Handling Missing Values:** We checked for NaN values in `df_modified`. In this specific dataset, no NaN values were found, so no replacement was necessary. If they were present, the strategy was to replace them with the overall mean of numerical values.
7. **Geographical Visualization:** A scatter plot of 'latitude' versus 'longitude' was generated to visualize the geographical distribution of the housing data points, revealing clusters and patterns in property locations.
8. **Min-Max Normalization:** We applied Min-Max normalization to 'median\_income' and 'housing\_median\_age' columns. This transformed the values to a scale between 0 and 1, which is crucial for many machine learning